

Lab 2 - Data wrangling

Sheila Yu

2026-02-04

```
library(tidyverse)
data <- midwest
```

Questions

Part 1

Question 1

```
counties_by_state <- data %>%
  count(state, name = "n_counties") %>%
  arrange(desc(n_counties))
```

```
counties_by_state
```

```
# A tibble: 5 x 2
  state n_counties
  <chr>     <int>
1 IL          102
2 IN           92
3 OH           88
4 MI           83
5 WI           72
```

The code above displays calculation the number of counties in each state. Based on the calculation, Illinois (IL) has the highest number of counties, with a total of 102 counties. Wisconsin (WI) has the lowest number of counties, with a total of 72 counties.

Question 2

```
data %>%
  count(county, name = "n_states") %>%
  filter(n_states == n_distinct(midwest$state))

# A tibble: 3 x 2
  county   n_states
  <chr>     <int>
1 CRAWFORD      5
2 JACKSON        5
3 MONROE         5
```

Based on the result above, the county that fills in the blank of “Look at that, there is a county called _____ in each state in this dataset!” includes Crawford, Jackson and Monroe.

Question 3

Part a:

```
data %>%
  filter(popdensity > 25000) %>%
  select(county, state, popdensity, poptotal, area) %>%
  arrange(desc(popdensity))
```

```
# A tibble: 9 x 5
  county    state popdensity poptotal   area
  <chr>     <chr>    <dbl>      <int> <dbl>
1 COOK      IL        88018.  5105067 0.058
2 MILWAUKEE WI        63952.  959275  0.015
3 WAYNE     MI        60334.  2111687 0.035
4 CUYAHOGA OH        54313.  1412140 0.026
5 DU PAGE   IL        39083.  781666  0.02
6 MARION    IN        34659.  797159  0.023
7 HAMILTON OH        34649.  866228  0.025
8 FRANKLIN OH        28278.  961437  0.034
9 MACOMB    MI        25621.  717400  0.028
```

Based on the results above, there are a total of nine counties with a population density above 25,000. These counties include: Cook, Milwaukee, Wayne, Cuyahoga, Du Page, Marion, Hamilton, Franklin and Macomb.

Part b:

```
data %>%
  filter(popdensity == max(popdensity, na.rm = TRUE)) %>%
  select(county, state, popdensity, poptotal, area)
```

```
# A tibble: 1 x 5
  county state popdensity poptotal   area
  <chr>  <chr>    <dbl>      <int> <dbl>
1 COOK    IL        88018.  5105067 0.058
```

Based on the result above, the county with the highest population density is Cook.

Question 4

```
data %>%
  summarise(
    median_density = median(popdensity, na.rm = TRUE),
    q1 = quantile(popdensity, 0.25, na.rm = TRUE),
    q3 = quantile(popdensity, 0.75, na.rm = TRUE)
  )

# A tibble: 1 x 3
  median_density     q1     q3
  <dbl>      <dbl>   <dbl>
1       1156.    622.   2330
```

The distribution of population density of counties is unimodal and extremely right-skewed. A typical Midwestern county has population density of 1156 people per unit area. The middle 50% of the counties have population densities between 622 to 2330 people per unit area.

*Please note that the data are rounded to the nearest whole number.

Question 5

```
data %>%
  mutate(metro = if_else(inmetro == 1, "Yes", "No")) %>%
  count(state, metro) %>%
  group_by(state) %>%
  mutate(proportion = n / sum(n)) %>%
  ungroup()
```

```
# A tibble: 10 x 4
  state metro      n proportion
  <chr> <chr> <int>     <dbl>
1 IL    No        74     0.725
2 IL    Yes       28     0.275
3 IN    No        55     0.598
4 IN    Yes       37     0.402
5 MI    No        58     0.699
6 MI    Yes       25     0.301
7 OH    No        48     0.545
8 OH    Yes       40     0.455
9 WI    No        52     0.722
10 WI   Yes      20     0.278
```

Here is a table with results from the calculation above displayed and rounded off to nearest percentage:

State	Metro or not	Proportion of population in the area
IL	No	73%
IL	Yes	27%
IN	No	60%
IN	Yes	40%
MI	No	70%
MI	Yes	30%
OH	No	55%
OH	Yes	45%
WI	No	72%
WI	Yes	28%

Question 6

Part a:

```
data %>%
  filter(
    percollege < 12,
    percbelowpoverty > 45
  ) %>%
  select(county, state, percbelowpoverty, percollege)
```

```
# A tibble: 1 x 4
  county     state percbelowpoverty percollege
  <chr>      <chr>        <dbl>       <dbl>
1 MENOMINEE WI            48.7        7.34
```

The outlier in the orange square is the county Menominee.

Part b:

```
data %>%
  filter(
    percollege > 40,
    percbelowpoverty < 20
  ) %>%
  select(county, state, percbelowpoverty, percollege)
```

```
# A tibble: 5 x 4
  county     state percbelowpoverty percollege
  <chr>      <chr>        <dbl>       <dbl>
1 CHAMPAIGN IL            15.6        41.3
2 DU PAGE   IL            2.71         42.8
3 HAMILTON IN            3.59         42.1
4 WASHTENAW MI           12.2         48.1
5 DANE       WI            10.5         43.6
```

The outliers in the red square are: Champaign, Du Page, Hamilton, Washtenaw, Dane.

Part c:

```

data %>%
  filter(
    (percollege < 12 & percbelowpoverty > 45) |
    (percollege > 40 & percbelowpoverty < 20)
  ) %>%
  select(county, state, percbelowpoverty, percollege)

```

```

# A tibble: 6 x 4
  county     state percbelowpoverty percollege
  <chr>      <chr>        <dbl>       <dbl>
1 CHAMPAIGN IL            15.6        41.3
2 DU PAGE   IL            2.71         42.8
3 HAMILTON IN            3.59         42.1
4 WASHTEENAW MI           12.2        48.1
5 DANE       WI            10.5        43.6
6 MENOMINEE WI           48.7         7.34

```

Above is a list of all the outliers in the orange and red square combined.

Part d:

```

midwest <- midwest %>%
  mutate(
    potential_outlier = if_else(
      (percollege < 12 & percbelowpoverty > 45) |
      (percollege > 40 & percbelowpoverty < 20),
      "Yes",
      "No"
    )
  )

midwest %>%
  select(county, state, percbelowpoverty, percollege, potential_outlier) %>%
  arrange(potential_outlier)

```

```

# A tibble: 437 x 5
  county     state percbelowpoverty percollege potential_outlier
  <chr>      <chr>        <dbl>       <dbl> <chr>
1 ADAMS      IL            13.2        19.6 No
2 ALEXANDER  IL            32.2        11.2 No
3 BOND       IL            12.1        17.0 No

```

```

4 BOONE    IL        7.21      17.3 No
5 BROWN    IL        13.5       14.5 No
6 BUREAU   IL        10.4       18.9 No
7 CALHOUN  IL        15.1       11.9 No
8 CARROLL  IL        11.7       16.2 No
9 CASS     IL        13.9       14.1 No
10 CHRISTIAN IL       11.7      13.6 No
# i 427 more rows

```

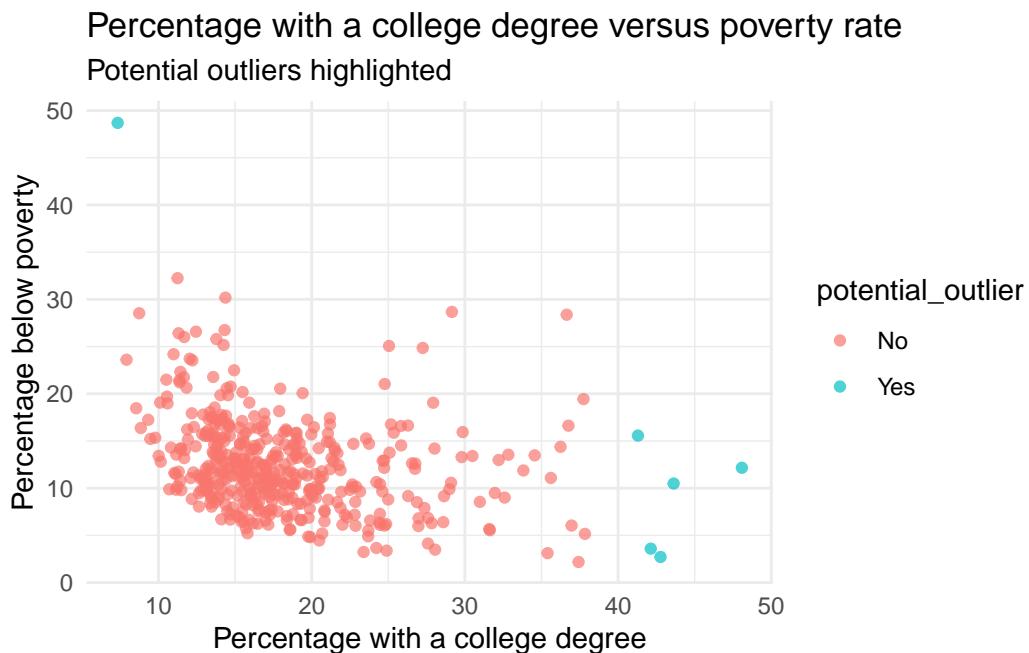
The updated dataset, `midwest`, now contains an additional column called `potential_outlier`, which indicates whether a county is included in the list of potential outliers identified in part c.

Part e:

```

ggplot(midwest, aes(x = percollege, y = percbelowpoverty, color = potential_outlier)) +
  geom_point(alpha = 0.7) +
  labs(
    x = "Percentage with a college degree",
    y = "Percentage below poverty",
    title = "Percentage with a college degree versus poverty rate",
    subtitle = "Potential outliers highlighted"
  ) +
  theme_minimal()

```



Above is the visualization of the relationship between college education and poverty, with colour highlighting the potential outliers in blue.

Question 7

Part a:

```
state_population <- midwest %>%
  group_by(state) %>%
  summarise(total_population = sum(poptotal, na.rm = TRUE)) %>%
  arrange(desc(total_population))

state_population
```

```
# A tibble: 5 x 2
  state total_population
  <chr>      <int>
1 IL          11430602
2 OH          10847115
3 MI          9295297
4 IN          5544159
5 WI          4891769
```

Part b:

```
state_population %>%
  mutate(
    population_proportion = total_population / sum(total_population)
  ) %>%
  arrange(desc(population_proportion))
```

```
# A tibble: 5 x 3
  state total_population population_proportion
  <chr>      <int>            <dbl>
1 IL          11430602        0.272
2 OH          10847115        0.258
3 MI          9295297         0.221
4 IN          5544159         0.132
5 WI          4891769         0.116
```

Part c: Illinois is the most populous Midwestern state, containing 27% of the total Midwest population. Wisconsin is the least populous state, accounting for only 12% of the Midwest population.

Question 8

```
state_poverty <- midwest %>%
  group_by(state) %>%
  summarise(
    mean_percbelowpoverty = mean(percbelowpoverty, na.rm = TRUE)
  )

state_poverty %>%
  arrange(mean_percbelowpoverty)
```

```
# A tibble: 5 x 2
  state mean_percbelowpoverty
  <chr>          <dbl>
1 IN              10.3
2 WI              11.9
3 OH              13.0
4 IL              13.1
5 MI              14.2
```

Based on the result above, Wisconsin has the lowest average percentage of population (10%) below the poverty line across its counties. Mississippi has the highest average percentage of population (14%) below the poverty line across its counties.

Part 2

Setting up data frame:

```
df <- tibble(  
  var_1 = c(10, 20, 30, 40, 50),  
  var_2 = c("Pizza", "Burger", "Pizza", "Pizza", "Burger"),  
  var_3 = c("Apple", "Apple", "Pear", "Pear", "Banana")  
)  
  
df
```

```
# A tibble: 5 x 3  
  var_1 var_2  var_3  
  <dbl> <chr>  <chr>  
1     10 Pizza  Apple  
2     20 Burger Apple  
3     30 Pizza  Pear  
4     40 Pizza  Pear  
5     50 Burger Banana
```

Question 9

Part a:

```
df |>  
  arrange(var_2)
```

```
# A tibble: 5 x 3  
  var_1 var_2  var_3  
  <dbl> <chr>  <chr>  
1     20 Burger Apple  
2     50 Burger Banana  
3     10 Pizza  Apple  
4     30 Pizza  Pear  
5     40 Pizza  Pear
```

The arrange () function reorders the rows of the data frame alphabetically by data in the column var_2.

Part b:

```
df |>  
  group_by(var_2)
```

```
# A tibble: 5 x 3  
# Groups: var_2 [2]  
  var_1 var_2 var_3  
  <dbl> <chr> <chr>  
1     10 Pizza Apple  
2     20 Burger Apple  
3     30 Pizza Pear  
4     40 Pizza Pear  
5     50 Burger Banana
```

The group_by () function groups the data frame by the variable var_2. It doesn't change the appearance of the data set, but it does createsgroups internally so that subsequent operations are performed separately for each group.

Part c:

```
df |>  
  group_by(var_2) |>  
  summarise(mean_var_1 = mean(var_1))
```

```
# A tibble: 2 x 2  
  var_2  mean_var_1  
  <chr>      <dbl>  
1 Burger      35  
2 Pizza       26.7
```

The pipeline first groups the data by var_2, then calculates the mean of var_1 within each group that is created. Therefore it computes the mean of var_1 separately for each category in var_2.

Part d:

```
df |>  
  group_by(var_2, var_3) |>  
  summarise(mean_var_1 = mean(var_1))
```

```
`summarise()` has grouped output by 'var_2'. You can override using the  
.groups` argument.
```

```

# A tibble: 4 x 3
# Groups:   var_2 [2]
  var_2  var_3  mean_var_1
  <chr>  <chr>     <dbl>
1 Burger Apple      20
2 Burger Banana     50
3 Pizza  Apple      10
4 Pizza  Pear       35

```

This pipeline groups the data frame by both var_2 and var_3, and then calculates the mean of var_1 within each unique combination of var_2 and var_3. The message displays:

```
`summarise()` has grouped output by 'var_2'. You can override using the `.groups` argument.
```

It indicates that after applying the summarize () function, the resulting data frame remains grouped by var_2. The message also notes that this behavior can be overridden by specifying the .groups argument, if an ungrouped output is desired.

Part e:

```

df |>
  group_by(var_2, var_3) |>
  summarise(mean_var_1 = mean(var_1), .groups = "drop")

```

```

# A tibble: 4 x 3
  var_2  var_3  mean_var_1
  <chr>  <chr>     <dbl>
1 Burger Apple      20
2 Burger Banana     50
3 Pizza  Apple      10
4 Pizza  Pear       35

```

This pipeline repeats the same grouped mean calculation as part d, but the .groups = “drop” argument tells summarize () to remove all grouping from the output, thereby un-grouping the output. Compared to the output in part d, the output in part 3 results in tibbles that is not grouped. No message (as shown in part d) is showing, suggesting that the output is a regular ungrouped data frame. However, it is also worth to note that the rows and values are the same in both part d and e.

Part f:

```
df |>
  group_by(var_2, var_3) |>
  summarize(mean_var_1 = mean(var_1), .groups = "drop")
```

```
# A tibble: 4 x 3
  var_2   var_3   mean_var_1
  <chr>   <chr>     <dbl>
1 Burger  Apple      20
2 Burger  Banana     50
3 Pizza   Apple      10
4 Pizza   Pear       35
```

```
df |>
  group_by(var_2, var_3) |>
  mutate(mean_var_1 = mean(var_1))
```

```
# A tibble: 5 x 4
# Groups:   var_2, var_3 [4]
  var_1 var_2   var_3   mean_var_1
  <dbl> <chr>   <chr>     <dbl>
1     10 Pizza   Apple      10
2     20 Burger  Apple      20
3     30 Pizza   Pear       35
4     40 Pizza   Pear       35
5     50 Burger  Banana     50
```

Pipeline 1 collapses the data and returns one row per group (var_2, var_3) with the mean of var_1. On the other hand, pipeline 2 does not collapse the data. Instead, it keeps all original rows, but adds a new column (mean_var_1) on top of existing rows where every row in the same group gets the same group mean value.